# State Data Revisted (OPTIONAL)

Bookmark this page

**IMPORTANT NOTE: This problem is optional, and will not count towards your grade. We have created this problem to give you extra practice with the topics covered in this unit.**

## State data Revisited (OPTIONAL)

We will be revisiting the "state" dataset from one of the optional problems in Unit 2. This dataset has, for each of the fifty U.S. states, the population, per capita income, illiteracy rate, murder rate, high school graduation rate, average number of frost days, area, latitude and longitude, division the state belongs to, region the state belongs to, and two-letter abbreviation. This dataset comes from the U.S. Department of Commerce, Bureau of the Census.

Load the dataset into R and convert it to a data frame by running the following two commands in R:

```
data(state)
statedata = data.frame(state.x77)
```

If you can't access the state dataset in R, here is a CSV file with the same data that you can load into R using the read.csv function: statedataSimple.csv. Be sure to call the output of the read.csv function "statedata".

After you have loaded the data into R, inspect the data set using the command: str(statedata)

This dataset has 50 observations (one for each US state) and the following 8 variables:

- **Population** - the population estimate of the state in 1975
- **Income** - per capita income in 1974
- **Illiteracy** - illiteracy rates in 1970, as a percent of the population
- **Life.Exp** - the life expectancy in years of residents of the state in 1970
- **Murder** - the murder and non-negligent manslaughter rate per 100,000 population in 1976
- **HS.Grad** - percent of high-school graduates in 1970
- **Frost** - the mean number of days with minimum temperature below freezing from 1931–1960 in the capital or a large city of the state
- **Area** - the land area (in square miles) of the state

We will try to build a model for life expectancy using regression trees, and employ cross-validation to improve our tree's performance.

## Problem 1.1 - Linear Regression Models

0 points possible (ungraded)

Let's recreate the **linear regression** models we made in the previous homework question. First, predict *Life.Exp* using all of the other variables as the independent variables (*Population, Income, Illiteracy, Murder, HS.Grad, Frost, Area* ). Use the entire dataset to build the model.

What is the **adjusted** R-squared of the model?

Answer: 0.6922

Explanation
To build the regression model, type the following command in your R console:
RegModel = lm(Life.Exp ~ ., data=statedata)
Then, if you look at the output of summary(RegModel), you should see that the Adjusted R-squared is 0.6922.

Submit

You have used 0 of 3 attempts

Calculator

Answers are displayed within the problem

## Problem 1.2 - Linear Regression Models

0 points possible (ungraded)
Calculate the sum of squared errors (SSE) between the predicted life expectancies using this model and the actual life expectancies:

Answer: 23.3

Explanation
To make predictions, type in your R console:
Predictions = predict(RegModel)
where "RegModel" is the name of your regression model. You can then compute the sum of squared errors by typing the following in your R console:
sum((statedata$Life.Exp - Predictions)^2)
The SSE is 23.29714.
Alternatively, you can use the following command to get the SSE:
sum(RegModel$residuals^2)

Submit You have used 0 of 3 attempts

Answers are displayed within the problem

## Problem 1.3 - Linear Regression Models

0 points possible (ungraded)
Build a second **linear regression** model using just *Population, Murder, Frost, and HS.Grad* as independent variables (the best 4 variable model from the previous homework). What is the **adjusted** R-squared for this model?

Answer: 0.71

Explanation
You can create this regression model by typing the following into your R console:
RegModel2 = lm(Life.Exp ~ Population + Murder + Frost + HS.Grad, data=statedata)
Then, if you type:
summary(RegModel2)
The Adjusted R-squared is at the bottom right of the output, and is 0.7126

Submit You have used 0 of 3 attempts

Answers are displayed within the problem

## Problem 1.4 - Linear Regression Models

0 points possible (ungraded)
Calculate the sum of squared errors again, using this reduced model:

Answer: 23.3

Explanation
The sum of squared errors (SSE) can be computed by first making predictions:
Predictions2 = predict(RegModel2)
and then computing the sum of the squared difference between the actual values and the predictions:
sum((statedata$Life.Exp - Predictions2)^2).

Calculator

Alternatively, you can compute the SSE with the following command:
SSE = sum(RegModel2$residuals^2)

Submit You have used 0 of 3 attempts

Answers are displayed within the problem

## Problem 1.5 - Linear Regression Models

0 points possible (ungraded)
Which of the following is correct?

- ✔ Trying different combinations of variables in linear regression is like trying different numbers of splits in a tree - this controls the complexity of the model.
- Using many variables in a linear regression is **always** better than using just a few.
- The variables we removed were uncorrelated with *Life.Exp*

Explanation
The correct answer is the first one. Trying different combinations of variables in linear regression controls the complexity of the model. This is similar to trying different numbers of splits in a tree, which is also controlling the complexity of the model.
The second answer is incorrect because as we see here, a model with fewer variables actually has a higher adjusted R-squared. If your accuracy is just as good, a model with fewer variables is almost always better.
The third answer is incorrect because the variables we removed have non-zero correlations with the dependent variable Life.Exp. Illiteracy and Area are negatively correlated with Life.Exp, with correlations of -0.59 and -0.11. Income is positively correlated with Life.Exp, with a correlation of 0.34. These correlations can be computed by typing the following into your R console:
cor(statedata$Life.Exp, statedata$Income)
cor(statedata$Life.Exp, statedata$Illiteracy)
cor(statedata$Life.Exp, statedata$Area)

Submit You have used 0 of 1 attempt

Answers are displayed within the problem

## Problem 2.1 - CART Models

0 points possible (ungraded)
Let's now build a **CART model** to predict *Life.Exp* using all of the other variables as independent variables (*Population, Income, Illiteracy, Murder, HS.Grad, Frost, Area*). We'll use the default *minbucket* parameter, so don't add the *minbucket* argument. Remember that in this problem we are not as interested in *predicting* life expectancies for new observations as we are understanding how they relate to the other variables we have, so we'll use all of the data to build our model. You shouldn't use the method="class" argument since this is a regression tree.

Plot the tree. Which of these variables appear in the tree? Select all that apply.

- [ ] Population
- [✔] Murder
- [ ] Frost

Calculator

- [ ] HS.Grad
- [ ] Area

Explanation
You can create the tree in R by typing the following command:
CARTmodel = rpart(Life.Exp ~ ., data=statedata)
Be sure to load the "rpart" and "rpart.plot" packages with the library command if they are not already loaded.
You can then plot the tree by typing:
prp(CARTmodel)
We can see that the only variable used in the tree is "Murder".

Submit You have used 0 of 3 attempts

Answers are displayed within the problem

## Problem 2.2 - CART Models

0 points possible (ungraded)
Use the regression tree you just built to predict life expectancies (using the predict function), and calculate the sum-of-squared-errors (SSE) like you did for linear regression. What is the SSE?

Answer: 29.0

Explanation
You can make predictions using the CART model by typing the following line in your R console (assuming your model is called "CARTmodel"):
PredictionsCART = predict(CARTmodel)
Then, you can compute the sum of squared errors (SSE) by typing the following:
sum((statedata$Life.Exp - PredictionsCART)^2)
The SSE is 28.99848.

Submit You have used 0 of 3 attempts

Answers are displayed within the problem

## Problem 2.3 - CART Models

0 points possible (ungraded)
The error is higher than for the linear regression models. One reason might be that we haven't made the tree big enough. Set the *minbucket* parameter to 5, and recreate the tree.

Which variables appear in this new tree? Select all that apply.

- [ ] Population
- [x] Murder
- [ ] Frost
- [x] HS.Grad

Calculator

- [x] Area

Explanation
You can create a tree with a minbucket value of 5 with the following command:
CARTmodel2 = rpart(Life.Exp ~ ., data=statedata, minbucket=5)
Then, if you plot the tree using prp(CARTmodel2), you can see that Murder, HS.Grad, and Area are all used in this new tree.

Submit You have used 0 of 3 attempts

Answers are displayed within the problem

## Problem 2.4 - CART Models

0 points possible (ungraded)
Do you think the default minbucket parameter is smaller or larger than 5 based on the tree that was built?

- [ ] Smaller
- [x] Larger

Explanation
Since the tree now has more splits, it must be true that the default minbucket parameter was limiting the tree from splitting more before. So the default minbucket parameter must be larger than 5.

Submit You have used 0 of 1 attempt

Answers are displayed within the problem

## Problem 2.5 - CART Models

0 points possible (ungraded)
What is the SSE of this tree?

Answer: 23.6

Explanation
You can compute the SSE of this tree by first making predictions:
PredictionsCART2 = predict(CARTmodel2)
and then computing the sum of the squared differences between the actual values and the predicted values:
sum((statedata$Life.Exp - PredictionsCART2)^2)
The SSE is 23.64283

This is much closer to the linear regression model's error. By changing the parameters we have improved the fit of our model.

Submit You have used 0 of 3 attempts

Answers are displayed within the problem

Calculator

## Problem 2.6 - CART Models

0 points possible (ungraded)

Can we do even better? Create a tree that predicts *Life.Exp* using **only** *Area*, with the *minbucket* parameter to 1. What is the SSE of this newest tree?

Answer: 9.3

Explanation

You can create this third tree by typing:

CARTmodel3 = rpart(Life.Exp ~ Area, data=statedata, minbucket=1)

Then to compute the SSE, first make predictions:

PredictionsCART3 = predict(CARTmodel3)

And then compute the sum of squared differences between the actual values and the predicted values:

sum((statedata$Life.Exp - PredictionsCART3)^2)

The SSE is 9.312442.

Note that the SSE is not zero here - we still make some mistakes. This is because there are other parameters in rpart that are also trying to prevent the tree from overfitting by setting default values. So our tree doesn't necessarily have one observation in each bucket - by setting minbucket=1 we are just allowing the tree to have one observation in each bucket.

Submit You have used 0 of 3 attempts

Answers are displayed within the problem

## Problem 2.7 - CART Models

0 points possible (ungraded)

This is the lowest error we have seen so far. What would be the best interpretation of this result?

- Trees are much better than linear regression for this problem because they can capture nonlinearities that linear regression misses.
- We can build almost perfect models given the right parameters, even if they violate our intuition of what a good model should be. ✔
- Area is obviously a very meaningful predictor of life expectancy, given we were able to get such low error using just Area as our independent variable.

Explanation

The correct answer is the second one. By making the minbucket parameter very small, we could build an almost perfect model using just one variable, that is not even our most significant variable. However, if you plot the tree using prp(CARTmodel3), you can see that the tree has 22 splits! This is not a very interpretable model, and will not generalize well.

The first answer is incorrect because our tree model that was not overfit performed similarly to the linear regression model. Trees only look better than linear regression here because we are overfitting the model to the data.

The third answer is incorrect because Area is not actually a very meaningful predictor. Without overfitting the tree, our model would not be very accurate only using Area.

Submit You have used 0 of 1 attempt

Answers are displayed within the problem

## Problem 3.1 - Cross-validation

Calculator

0 points possible (ungraded)
Adjusting the variables included in a linear regression model is a form of model tuning. In Problem 1 we showed that by removing variables in our linear regression model (tuning the model), we were able to maintain the fit of the model while using a simpler model. A rule of thumb is that simpler models are more interpretable and generalizeable. We will now tune our regression tree to see if we can improve the fit of our tree while keeping it as simple as possible.

Load the *caret* library, and set the seed to 111. Set up the controls exactly like we did in the lecture (10-fold cross-validation) with *cp* varying over the range 0.01 to 0.50 in increments of 0.01. Use the *train* function to determine the best *cp* value for a CART model using all of the available independent variables, and the entire dataset statedata. What value of cp does the train function recommend? (Remember that the train function tells you to pick the largest value of cp with the lowest error when there are ties, and explains this at the bottom of the output.)

Answer: 0.12

Explanation
You can load the library caret and set the seed by typing the following commands:
library(caret)
set.seed(111)
Then, you can set up the cross-validation controls by typing:
fitControl = trainControl(method = "cv", number = 10)
cartGrid = expand.grid(.cp = seq(0.01, 0.5, 0.01) )
You can then use the train function to find the best value of cp by typing:
train(Life.Exp ~ ., data=statedata, method="rpart", trControl = fitControl, tuneGrid = cartGrid)
At the bottom of the output, it says that the best value of cp is 0.12.

Submit You have used 0 of 4 attempts

Answers are displayed within the problem

## Problem 3.2 - Cross-Validation

0 points possible (ungraded)
Create a tree with the value of *cp* you found in the previous problem, all of the available independent variables, and the entire dataset "statedata" as the training data. Then plot the tree. You'll notice that this is actually quite similar to the first tree we created with the initial model. Interpret the tree: we predict the life expectancy to be 70 if the murder rate is greater than or equal to

Answer: 6.6

and is less than

Answer: 11

Explanation
You can create a new tree with cp=0.12 by typing:
CARTmodel4 = rpart(Life.Exp ~ ., data=statedata, cp=0.12)
Then, if you plot the tree using prp(CARTmodel4), you can see that the life expectancy is predicted to be 70 if Murder is greater than or equal to 6.6 (the first split) and less than 11 (the second split).

Submit You have used 0 of 4 attempts

Answers are displayed within the problem

## Problem 3.3 - Cross-Validation

Calculator

0 points possible (ungraded)
Calculate the SSE of this tree:

Answer: 32.9

Explanation
To compute the SSE, first make predictions:
PredictionsCART4 = predict(CARTmodel4)
and then compute the sum of squared differences between the actual values and the predicted values:
sum((statedata$Life.Exp - PredictionsCART4)^2)
The SSE is 32.86549

Submit You have used 0 of 3 attempts

Answers are displayed within the problem

## Problem 3.4 - Cross-Validation

0 points possible (ungraded)
Recall the first tree (default parameters), second tree (minbucket = 5), and the third tree (selected with cross validation) we made. Given what you have learned about cross-validation, which of the three models would you expect to be better if we did use it for prediction on a test set? For this question, suppose we had actually set aside a few observations (states) in a test set, and we want to make predictions on those states.

- The first model
- The second model
- The model we just made with the "best" cp ✔

Explanation
The purpose of cross-validation is to pick the tree that will perform the best on a test set. So we would expect the model we made with the "best" cp to perform best on a test set.

Submit You have used 0 of 1 attempt

Answers are displayed within the problem

## Problem 3.5 - Cross-Validation

0 points possible (ungraded)
At the end of Problem 2 we made a very complex tree using just Area. Use *train* with the same parameters as before but just using Area as an independent variable to find the best cp value (set the seed to 111 first). Then build a new tree using just Area and this value of cp.

How many splits does the tree have?

Answer: 4

Explanation
To find the best value of cp when using only Area, use the following command:
set.seed(111)
train(Life.Exp ~ Area, data=statedata, method="rpart", trControl = fitControl, tuneGrid = cartGrid )
Then, build a new CART tree by typing:

Calculator

Then, build a new CART tree by typing:
CARTmodel5 = rpart(Life.Exp ~ Area, data=statedata, cp=0.02)
You can plot the tree with prp(CARTmodel5), and see that the tree has 4 splits.

Submit You have used 0 of 4 attempts

Answers are displayed within the problem

## Problem 3.6 - Cross-Validation

0 points possible (ungraded)
The lower left leaf (or bucket) corresponds to the lowest predicted Life.Exp of 70. Observations in this leaf correspond to states with area greater than or equal to

Answer: 9579

and area less than

Answer: 51000

Explanation
To get to this leaf, we go through 3 splits:
Area less than 62,000
Area greater than or equal to 9,579
Area less than 51,000
This means that this leaf is composed of states that have an area greater than 9,579 and less than 51,000.

Submit You have used 0 of 4 attempts

Answers are displayed within the problem

## Problem 3.7 - Cross-Validation

0 points possible (ungraded)
We have simplified the previous "Area tree" considerably by using cross-validation. Calculate the SSE of the cross-validated "Area tree", and select all of the following correct statements that apply:

- [ ] The best model in this whole question is the first "Area tree" because it had the lowest SSE.
- [x] The Area variable is not as predictive as Murder rate.
- [ ] Cross-validation is intended to decrease the SSE for a model on the training data, compared to a tree that isn't cross-validated.
- [ ] Cross-validation will always improve the SSE of a model on unseen data, compared to a tree that isn't cross-validated.

Explanation
You can calculate the SSE by first making predictions and then computing the SSE:
PredictionsCART5 = predict(CARTmodel5)
sum((statedata$Life.Exp - PredictionsCART5)^2)
The original Area tree was overfitting the data - it was uninterpretable. Area is not as useful as Murder - if it was, it would have been in the cross-validated tree. Cross-validation is not designed to improve the fit on the training data, but it won't necessarily make it worse either. Cross-validation cannot guarantee improving the SSE on unseen data, although it often helps.

Calculator

SSE on unseen data, although it often helps.

Submit You have used 0 of 2 attempts

Answers are displayed within the problem

Please remember not to ask for or post complete answers to homework questions in this discussion forum.

## Discussion

**Topic:** Unit 4 / Unit 4, Homework: State Data Revisted

**Show Discussion**

‹ Previous Next ›

## Connect

Idea Hub
Contact Us
Help Center
Security
Media Kit

GET IT ON Google play
Download on the App Store


深圳市恒宇博科技有限公司 粤ICP备17044299号-2

Calculator